



SCS492-Selected Topics in Software Engineering-2 [Machine Learning]
Assignment - 1

Supervised By:
DR.Hanaa Mobarez
TA.Abdalrahman Roshdi

Prepared & Implemented By:
Mootaz Medhat Ezzat Abdelwahab (20206074)

➤ Documentation of Code Execution and Results:

- a) Load the "loan_old.csv" dataset
- h) Load "loan_new.csv" dataset
- b) Perform analysis on the dataset:
 - i) check whether there are missing values in "loan_old.csv" dataset
check whether there are missing values in "loan_new.csv" dataset
 - ii) check the type of each column (categorical or numerical)
 - iii) check whether numerical columns have the same scale
 - iv) visualize a pairplot between numerical columns
- c) Preprocess the data:
 - i) remove records with missing values from the original "loan_old" DataFrame directly and returns None
remove records with missing values from the original "loan_new" DataFrame directly and returns None
 - ii) separate features and targets
 - iii) shuffle and split the data into training and testing sets using train_test_split() function from scikit-learn, it randomly shuffles the data before splitting it into training and testing sets
 - iv) encode categorical features into numerical labels
 - v) encode categorical targets into numerical labels
 - vi) standardize the features using the mean and standard deviation
- d) Fit linear regression model
- e) Evaluate linear regression model
- f) Fit logistic regression model from scratch
- g) Function (from scratch) to calculate the accuracy of the logistic regression model
- j) Predict using models
- Print predictions for loan amounts and loan status for new data

➤ a) Load the "loan_old.csv" dataset:

```
11 # a) Load the "loan_old.csv" dataset:
12 # -----
13 loan_old = pd.read_csv("loan_old.csv")
```

➤ h) Load "loan_new.csv" dataset:

```
15 # h) Load "loan_new.csv" dataset:
16 # -----
17 loan_new = pd.read_csv("loan_new.csv")
```

➤ b) Perform analysis on the dataset:

• i) check whether there are missing values in "loan_old.csv" dataset

```
19 # b) Perform analysis on the dataset:
20 # -----
21 # i) check whether there are missing values in "loan_old.csv" dataset
22 print("Missing Values in loan_old Dataset:")
23 print(loan_old.isnull().sum(), "\n")
```

```
C:\Users\MAS\Documents\pyCharm\Machine Learning Assignments\Assignment_1\linear_and_logistic_regression\.venv\Scripts\python.exe "C:\Users\MAS\Documents\p
Missing Values in loan_old Dataset:
Loan_ID      0
Gender       13
Married       3
Dependents    15
Education     0
Income        0
Coapplicant_Income  0
Loan_Tenor    15
Credit_History  50
Property_Area  0
Max_Loan_Amount  25
Loan_Status   0
dtype: int64
```

• i) check whether there are missing values in "loan_new.csv" dataset

```
24 # check whether there are missing values in "loan_new.csv" dataset
25 print("Missing Values in loan_new Dataset:")
26 print(loan_new.isnull().sum(), "\n")
```

```
Missing Values in loan_new Dataset:
Loan_ID      0
Gender       11
Married       0
Dependents    10
Education     0
Income        0
Coapplicant_Income  0
Loan_Tenor     7
Credit_History  29
Property_Area  0
dtype: int64
```

- ii) check the type of each column (categorical or numerical)

```

28 # ii) check the type of each column (categorical or numerical)
29 print("Data Types of Each Column in loan_old Dataset:")
30 print("[float64, int64] means numerical")
31 print("[object] means categorical")
32 print("-----")
33 print(loan_old.dtypes, "\n")

```

```

Data Types of Each Column in loan_old Dataset:
[float64, int64] means numerical
[object] means categorical
-----
Loan_ID      object
Gender       object
Married      object
Dependents   object
Education    object
Income       int64
Coapplicant_Income float64
Loan_Tenor   float64
Credit_History float64
Property_Area object
Max_Loan_Amount float64
Loan_Status  object
dtype: object

```

- iii) check whether numerical columns have the same scale

```

35 # iii) check whether numerical columns have the same scale
36 print("Summary Statistics for the Numerical Features:")
37 print(loan_old.describe(), "\n")

```

```


Summary Statistics for the Numerical Features:
      Income  Coapplicant_Income  ...  Credit_History  Max_Loan_Amount
count  614.000000      614.000000  ...      564.000000      589.000000
mean   5403.459283      1621.245798  ...      0.842199      230.499474
std    6109.041673      2926.248369  ...      0.364878      161.976967
min     150.000000      0.000000    ...      0.000000      12.830000
25%    2877.500000      0.000000    ...      1.000000      123.990000
50%    3812.500000      1188.500000  ...      1.000000      190.370000
75%    5795.000000      2297.250000  ...      1.000000      276.500000
max     81000.000000     41667.000000  ...      1.000000      990.490000

[8 rows x 5 columns]

R^2 Score of the Linear Regression Model: 0.7970174519748384
Accuracy of the Logistic Regression Model: 0.8155339805825242

```

- iv) visualize a pairplot between numerical columns



➤ c) Preprocess the data:

- i) remove records with missing values from the original "loan_old" DataFrame directly and returns None

```
43 # c) Preprocess the data:
44 # -----
45 # i) remove records with missing values from the original "loan_old" DataFrame directly and returns None
46 loan_old.dropna(inplace=True)
```

- i) remove records with missing values from the original "loan_new" DataFrame directly and returns None

```
47 # remove records with missing values from the original "loan_new" DataFrame directly and returns None
48 loan_new.dropna(inplace=True)
```

- ii) separate features and targets

```
50 # ii) separate features and targets
51 X = loan_old.drop(columns=['Loan_ID', 'Max_Loan_Amount', 'Loan_Status'])
52 X_new = loan_new.drop(columns=['Loan_ID'])
53 y_amount = loan_old['Max_Loan_Amount']
54 y_status = loan_old['Loan_Status']
```

- iii) shuffle and split the data into training and testing sets using train_test_split() function from scikit-learn, it randomly shuffles the data before splitting it into training and testing sets

```
56 # iii) shuffle and split the data into training and testing sets using train_test_split() function from scikit-learn
57 # it randomly shuffles the data before splitting it into training and testing sets
58 X, y_amount, y_status = shuffle("arrays: X, y_amount, y_status")
59 X_train, X_test, y_amount_train, y_amount_test, y_status_train, y_status_test = train_test_split("arrays: X, y_amount, y_status", test_size=0.2)
60
61 X_train_encoded = X_train.copy()
62 X_test_encoded = X_test.copy()
63 X_new_encoded = X_new.copy()
64
65 categorical_columns = X.select_dtypes(include=['object']).columns
66 numerical_columns = X.select_dtypes(include=['int64', 'float64']).columns
```

- iv) encode categorical features into numerical labels

```
68 # iv) encode categorical features into numerical labels
69 encoder = LabelEncoder()
70 for column in categorical_columns:
71     X_train_encoded[column] = encoder.fit_transform(X_train_encoded[column])
72     X_test_encoded[column] = encoder.transform(X_test_encoded[column])
73     X_new_encoded[column] = encoder.transform(X_new_encoded[column])
```

- v) encode categorical targets into numerical labels

```
75 # v) encode categorical targets into numerical labels
76 y_status_train = encoder.fit_transform(y_status_train)
77 y_status_test = encoder.transform(y_status_test)
```

- vi) standardize the features using the mean and standard deviation

```
79 # vi) standardize the features using the mean and standard deviation
80 mean_values = X_train_encoded[numerical_columns].mean()
81 std_values = X_train_encoded[numerical_columns].std()
82
83 X_train_encoded[numerical_columns] = (X_train_encoded[numerical_columns] - mean_values) / std_values
84 X_test_encoded[numerical_columns] = (X_test_encoded[numerical_columns] - mean_values) / std_values
85 X_new_encoded[numerical_columns] = (X_new_encoded[numerical_columns] - mean_values) / std_values
```

➤ d) Fit linear regression model:

```
87 # d) Fit linear regression model:
88 # -----
89 linear_regression = LinearRegression()
90 linear_regression.fit(X_train_encoded, y_amount_train)
```

➤ e) Evaluate linear regression model:

```
92 # e) Evaluate linear regression model:
93 # -----
94 y_amount_predicted = linear_regression.predict(X_test_encoded)
95 r2 = r2_score(y_amount_test, y_amount_predicted)
96 print("R^2 Score of the Linear Regression Model: ", r2)
```

➤ f) Fit logistic regression model from scratch:

```
100 def sigmoid(z):
101     return 1 / (1 + np.exp(-z))
102
103     1 usage
104 def fit_GD(X, y):
105     m, n = X.shape
106
107     # initialize theta array with ones and size n+1 (theta.T)
108     theta = np.ones(n + 1)
109     # add a column of ones to X
110     X = np.column_stack((np.ones(m), X))
111
112     # set alpha and max_iterations
113     alpha = 0.01
114     # set max_iterations
115     max_iterations = 1000
116
117     # iterate over max_iterations
118     for i in range(max_iterations):
119         h = sigmoid(np.dot(X, theta))
120         # compute the partial derivative of the error:  $\partial J(\theta) / \partial \theta_j = (1/m) * X.T * (h - y)$ 
121         partial_derivative = (1 / m) * np.dot(X.T, (h - y))
122         # update theta_j according to the equation:  $\theta = \theta - \alpha * \partial J(\theta) / \partial \theta$ 
123         theta -= alpha * partial_derivative
124
125     return theta
```

```
134 def predict(X, theta):
135     return sigmoid(np.dot(X, theta))
136
137 X_train_logistic = X_train_encoded.copy()
138 theta = fit_GD(X_train_logistic.values, y_status_train)
```

➤ g) Function (from scratch) to calculate the accuracy of the logistic regression model:

```
140 # g) Function (from scratch) to calculate the accuracy of the logistic regression model:
141 # -----
142 1 usage
143 def accuracy(actual, predicted):
144     total_samples = len(actual)
145     correct_predictions = 0
146     for i in range(total_samples):
147         if predicted[i] >= 0.5:
148             if actual[i] == 1:
149                 correct_predictions += 1
150             else:
151                 if actual[i] == 0:
152                     correct_predictions += 1
153         return correct_predictions / total_samples
154
155 # Calculate the accuracy of the model
156 X_test_logistic = X_test_encoded.copy()
157 X_test_logistic = np.column_stack((np.ones(X_test_logistic.shape[0]), X_test_logistic))
158 y_status_pred = predict(X_test_logistic, theta)
159 model_accuracy = accuracy(y_status_test, y_status_pred)
160 print("Accuracy of the Logistic Regression Model: ", model_accuracy, "\n")
```

➤ j) Predict using models:

```
161 # j) Predict using models:
162 # -----
163 y_amount_new = linear_regression.predict(X_new_encoded)
164 X_predict_logistic = X_new_encoded.copy()
165 X_predict_logistic = np.column_stack((np.ones(X_predict_logistic.shape[0]), X_predict_logistic))
166 y_status_new = predict(X_predict_logistic, theta)
```

➤ Print predictions for loan amounts and loan status for new data

```
168 # Print predictions for loan amounts and loan status for new data
169 print("Predicted loan details for new data:")
170 print("-----")
171 for idx, (loan_id, amount, status) in enumerate(zip(loan_new['Loan_ID'], y_amount_new, y_status_new)):
172     print(f" - Loan ID: {loan_id}")
173     print(f" - Predicted Loan Amount: {amount:.2f}$")
174     print(f" - Predicted Loan Status: {'Approved (Y)' if status >= 0.5 else 'Rejected (N)'}")
175     print("-----")
```

```
- Loan ID: LP002747
- Predicted Loan Amount: 321.04$
- Predicted Loan Status: Rejected (N)
-----
- Loan ID: LP002759
- Predicted Loan Amount: 219.56$
- Predicted Loan Status: Approved (Y)
-----
- Loan ID: LP002760
- Predicted Loan Amount: 90.13$
- Predicted Loan Status: Approved (Y)
-----
- Loan ID: LP002766
- Predicted Loan Amount: 131.57$
- Predicted Loan Status: Approved (Y)
-----
- Loan ID: LP002769
- Predicted Loan Amount: 176.95$
- Predicted Loan Status: Approved (Y)
-----
- Loan ID: LP002774
- Predicted Loan Amount: 165.49$
- Predicted Loan Status: Rejected (N)
-----
```